


SYBILION HACKATHON 2026 · FORECASTING AI TRACK

Infineon Cost & Treasury Agent

Forecast → Hedge → Stock Safety

A decision agent on the Sybillion probabilistic-forecasting API that scopes a chip maker's USD-priced cost exposures — FX, wholesale power and commodities — forecasts each, keeps only the high-confidence signals, and turns them into a hedge or a safety-stock buffer, every recommendation walk-forward backtested on real data.

6-MONTH FORECAST · P10 / P50 / P90



TEAM

43 Food

Fridolin Sitter · Elliot Sezalory

Exposed on Three Fronts at Once

A European power-semiconductor maker like Infineon buys in USD but reports in EUR — and its margin is hostage to inputs it doesn't control. A single point forecast can't answer the treasury question:

THE TREASURY QUESTION

Which exposure is worth acting on, with **what conviction**, taking **what action** — and would acting on it **historically have helped**?

Why it matters: the interesting problem is **multi-exposure** — the same forecast bands must drive three different decisions, and the agent has to decide which and when, including when the honest answer is *“bands too wide, don't act.”*

WHERE THE MARGIN LEAKS



FX

USD-denominated supplier payments vs a EUR book



Wholesale power


Energy-intensive fabs exposed to day-ahead prices



Commodities

Copper, gallium and gas feed the cost basis

Four Layers, From Prompt to Proven Action



ONE SCORE GATES EVERYTHING $\text{conf} = 100 \cdot \tanh(\text{move} / \text{band}) \cdot (1 - \text{mape}) \rightarrow \text{STRONG only when the move beats the band}$

Designing for Uncertainty, Not Around It

ONE CONFIDENCE SCORE, *three decisions*

A signal is **STRONG** only when the predicted move is large relative to the forecast's own uncertainty. The same band width also sizes the safety-stock buffer.

```
band  = (p90 - p10) / p50
move  = |p50 - now| / now
conviction = move / band
conf  = 100 * tanh(conviction)
      · (1 - min(mape, 50) / 50)
```

Wide bands → more buffer weeks. **Directional bias** →
pre-buy vs run lean.



Forecast as a filter

For power & commodities we don't trade the forecast — we use Sybilion's direction to filter a Donchian breakout, then measure filtered vs unfiltered.



Walk-forward, never in-sample

Every backtest truncates history and scores against the realized next value, so accuracy is genuinely out-of-sample.



Mock-first architecture

Mock and live clients share one shape: the app runs offline with zero keys and swaps to live Sybilion with a single flag.



Shock scenarios

boj_hike · ecb_cut · china_deval re-bias the forecasts mid-run, so signals and actions visibly flip on stage.

We Report the Mock-to-Live Gap Honestly

83%

Mock forecasts

6 signals · 2026-05 · STRONG-only 2 / 2

40%

Live Sybilion

30 signals · 2026-01 → 2026-05

WHAT THE NUMBER MEANS

The polished demo (83%) and live reality (40%) are both shown — no cherry-picking. STRONG signals still beat WEAK in the mock set, and the confidence gate keeps the agent on the sidelines when the bands are too wide to justify a trade. Currency direction at a 1-month horizon is simply hard; the value is the discipline around the forecast, not a magic hit-rate.

THE HONEST READ

Monthly univariate forecasts are a **thin directional edge**. Across the live FX set we hit **40%** — roughly coin-flip at the monthly horizon for currencies.

Why the gap? We submit with `filters.limit = 0` — pure univariate, no external driver signals. That leaves Sybilion's most decision-relevant output on the table.


Artifacts: `monthly_backtest_live.json` · `pairs28_h6_live.json`

On Power & Commodities, the Filter Clearly Helps

Asset	Forecast	Baseline Donchian	+ Sybillion filter	Dir.
Copper · 132m	DOWN	10 trades · +0.9%	8 trades · -0.0% · ½ drawdown	20%
TTF gas · 132m	UP	11 trades · -2.5%	4 trades · -0.2% · cut losers	60%
AT power · 137m	UP	10 trades · -0.4%	3 trades · +0.3% · flipped +	60%

The win: the direction filter cuts trade count and drawdown on every asset, and on Austrian power it flips a losing strategy positive (-0.4% → +0.3%). The forecast earns its place as a *filter* — not as something to trade directly.

The honest caveat: none of these beat naive buy-and-hold over the long windows (copper +154%, power +254%). But “buy and hold a cost input forever” isn't a treasury action — the right comparison is filtered vs unfiltered, where the forecast clearly helps.



Band-width → safety stock

For materials we don't trade at all — the forecast's uncertainty sizes the buffer directly: wider bands buy more weeks of safety stock, a directional bias means pre-buy vs run lean.

The Clearest Path Past 40%



PRIORITY Turn on external drivers

Submit with `filters.limit > 0` and surface Sybillion's driver-importance per horizon — copper ← China export controls, power ← gas + carbon, FX ← rate differentials. Directly serves “visible reasoning,” and the likeliest path to beat 40%.



Cross-exposure synthesis

A USD spike raises both the FX-hedge need and material cost. Let the agent reason across exposures at once instead of one at a time.



Confidence-scaled sizing

Make hedge ratio and buffer weeks proportional to the score, plus a calibration report: are 70%-confidence signals right 70% of the time?



Longer live track record

Extend beyond 5 months and put transaction costs into the P&L for an honest, deployable read.